

Assignment 1

COMP7607: Natural Language Processing - The University of Hong Kong

Fall 2024

Due: October 31, 6:00 AM

Recent studies reveal that large language models (LLMs), such as Llama (Dubey et al., 2024) and GPT-4 (OpenAI, 2023), excel in step-by-step reasoning for addressing complex tasks driven by natural language instructions. In this assignment, you will explore the capabilities of the **Llama-3.1-8B-Instruct** model in two reasoning domains: **mathematics** and **coding**. You can choose one task—either mathematics or coding—based on your interests.

Submit: You should submit your assignment to the COMP7607 Moodle page. You will need to submit a PDF file **UniversityNumber.pdf** of your report (with **detailed** experimental details) and a zip file **UniversityNumber.zip**, which includes:

- .py files.
- zeroshot.baseline.jsonl
- fewshot.baseline.jsonl (for math task)
- [method_name_1].jsonl
- [method_name_2].jsonl
- [method_combine].jsonl or [method_improve].jsonl

Please note that the UniversityNumber is the number printed on your student card.

1 Introduction

Prompt Engineering refers to methods for how to instruct LLMs for desired outcomes without updating model weights. The core task of assignment 1 is to design methods for prompting LLMs to improve the accuracy of LLMs on math problem-solving or code generation.

Generally, we have two prompting paradigms:

1. **Zero-shot prompting** is to simply feed the task text to the model and ask for results. For example, we can send the following message to **Llama-3.1-8B-Instruct**:

```
messages = [
  {"role": "system", "content": [System Message]},
  {"role": "user", "content": "Elsa has 5 apples. Anna has 2 more apples than Elsa.
  How many apples do they have together?"},
]
```

“{"role": "system", "content": [System Message]}” can be optional. [System Message] is a feature-specific contextual description given to a generative AI model, such as “Your task is to solve a series of math word problems by providing the final answer. Use the format #### [value] to highlight your answer.”

2. **Few-shot prompting** presents a set of high-quality demonstrations, each consisting of both input and desired output, on the target task. As the model first sees good examples, it can better understand human intention and criteria for what kinds of answers are wanted. For example, we can send the following 1-shot prompting message to **Llama-3.1-8B-Instruct**:

```
messages = [
  {"role": "system", "content": [System Message]},
  {"role": "user", "content": "Elsa has 5 apples. Anna has 2 more apples than Elsa. How many apples do they have together?"},
  {"role": "assistant", "content": "Anna has 2 more apples than Elsa. So Anna has 2 + 5 = 7 apples. Elsa and Anna have 5 + 7 = 12 apples together. #### 12" },
  {"role": "user", "content": "If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?"},
]
```

The choice of prompting format, interaction method, and decoding parameters will all significantly influence task performance. Now enjoy exploring the magic of prompting with Llama-3.1-8B-Instruct!

2 Task

For your selected task (math or code), begin by implementing the two simplest baselines: zero-shot and few-shot prompting, using only the problem statement with demonstrations (optional). Please refer to the example format in Section 1. Then, choose two advanced methods to enhance performance beyond baselines. We present advanced methods that you may choose to implement in Section 2.1 and Section 2.2, respectively. You are also encouraged to propose your own designs!

Here, we provide two surveys that offer a comprehensive overview of the development of prompting engineering and the self-correction of LLMs:

- A survey for prompt engineering for autoregressive language models [[Survey Link](#)].
- A collection of research papers for LLMs self-correcting [[Survey Link](#)]. Note that we only focus on the “Generation-Time Correction” and “Post-hoc Correction”.

2.1 Mathematical Reasoning

Data We use the grade school math dataset, [GSM8K](#), which is one of the most popular datasets for evaluating the mathematical reasoning performance of LLMs. Note that the performance of the LLM is measured using “GSM8K/test.jsonl”. You can utilize “GSM8K/train.jsonl” for other objectives, such as retrieving the most similar examples for the questions in the test file.

Metric For each method, report the overall accuracy on the 1,319 questions in the test file. Use the answer extraction function provided in “GSM8K/evaluation.py” for calculating accuracy. Additionally, report the inference cost using two metrics: wall-clock time and the average number of generated tokens per math question.

Task 1: Implement Baseline We consider two baselines:

- Zero-shot prompt the model using only the problem statement.
- Few-shot prompts the model with a few demonstrations, each containing manually written (or model-generated) high-quality reasoning chains.

We provide the prompting template in “GSM8K/baseline.py”.

Task 2: Implement two Advanced Prompting Methods You may explore various advanced strategies for improvement, such as retrieving similar demonstrations, decomposing responses, employing response ensemble, engaging in self-improvement, and more. Below, we present several advanced methods, although not exhaustive, that you may choose to implement, considering the characteristics of mathematical reasoning tasks. You are also encouraged to propose your own designs!

- Self-Polish: Enhance Reasoning in Large Language Models via Problem Refinement [\[paper\]](#)
- Tree Prompting: Efficient Task Adaptation without Fine-Tuning [\[paper\]](#)
- SELF-REFINE: Iterative Refinement with Self-Feedback [\[paper\]](#)
- PHP: Progressive-Hint Prompting Improves Reasoning in Large Language Models [\[paper\]](#)
- CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing [\[paper\]](#)
- Teaching Algorithmic Reasoning via In-context Learning [\[paper\]](#)
- Contrastive Decoding Improves Reasoning in Large Language Models [paper\]](#)
- Self-Evaluation Guided Beam Search for Reasoning [\[paper\]](#)
- Skills-in-Context Prompting: Unlocking Compositionality in Large Language Models [\[paper\]](#)
- Program of Thoughts Prompting: Disentangling Computation from Reasoning for Numerical Reasoning Tasks [\[paper\]](#)
- Large Language Models are Better Reasoners with Self-Verification [\[paper\]](#)

Task 3: Combine the Two Methods or Enhance One Method Can you combine the two advanced methods or enhance one method based on your analysis to achieve greater gains? Are complementary or enhanced effects achievable? If not, please explain why.

2.2 Code Generation

A particularly intriguing application of LLMs is code generation, a task that involves producing source code from natural language descriptions. This area has garnered substantial interest from both academia and industry, as evidenced by the development of tools like GitHub Copilot¹.

Data We use [HumanEval leaderboard](#), which has been established as a de facto standard for evaluating the coding proficiency of LLMs. Note that the performance of the LLM is measured using “HumanEval/HumanEval.jsonl”.

Metric For each method, report the overall accuracy on the 164 questions in the “HumanEval.jsonl” file. We provided “example_problem.jsonl” and “example_solutions.jsonl” under HumanEval to illustrate the format and assist with debugging. You can evaluate the accuracy of the generated code by running:

```
python evaluate_functional_correctness.py -sample_file example_samples.jsonl -problem_file example_problem.jsonl
```

Also, report the inference cost using two metrics: wall-clock time and the average number of generated tokens per code problem.

Task 1: Implement Baseline We use zero-shot prompting as the baseline. The prompting template is provided in “HumanEval/baseline.py”.

¹<https://github.com/features/copilot>

Task 2: Implement two Advanced Prompting Methods You may explore various advanced strategies for improvement. Below, we present several advanced methods that you may choose to implement. You are also encouraged to propose your own designs!

- Teaching Large Language Models to Self-Debug [paper]
- SelfEvolve: A Code Evolution Framework via Large Language Models [paper]
- Self-Edit: Fault-Aware Code Editor for Code Generation [paper]
- Lever: Learning to verify language-to-code generation with execution [paper]
- Is Self-Repair a Silver Bullet for Code Generation? [paper]
- Codet: Code generation with generated tests [paper]
- SELF-REFINE: Iterative Refinement with Self-Feedback [paper]
- Language agent tree search unifies reasoning acting and planning in language models [paper]

Task 3: Combine the Two Methods or Enhance One Method Can you combine the two advanced methods or enhance one method based on your analysis to achieve greater gains? Are complementary or enhanced effects achievable? If not, please explain why.

Task 4: Report the Inference Cost

2.3 Model

In this assignment, we will use Llama-3.1-8B-Instruct, which is a powerful open-source model that natively supports multilingual capabilities, coding, reasoning, and tool usage. For more details about this model, you can refer to the Meta Blog: <https://ai.meta.com/blog/meta-llama-3-1/> and https://www.llama.com/docs/model-cards-and-prompt-formats/llama3_1/.

3 API

In this assignment, you will interact with the Llama-3.1-8B-Instruct sponsored by SambaNova System. To access this resource, please refer to the instructions in the “SambaNova Cloud QuickStart Guide.pdf” to register and generate your API key. To verify that your API key is functioning correctly, you can either use the provided curl command in the document or run the “test_full_response.py” script.

4 Report

You will write a report including the following parts:

- The description of your implemented methods, including the architecture, the hyperparameters, etc.
- The **discussion** and **table** of your results, including the hyperparameters you have tested, the performance of the methods on the test set, the analysis of the results, the advantages compared to the baselines, etc.

5 Note

There are some key points you should pay attention to:

- Your assignment will not be evaluated based solely on the accuracy of the task. Instead, we primarily assess your work through your experiments and analysis.
- You can adjust the decoding hyperparameters to get a better performance on the test set. The hyperparameters include the temperature, top-p, top-k, etc.
- The Llama-3.1-8B-Instruct is the only base model to be used for solving tasks in this assignment.

References

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurélien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Rozière, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Graeme Nail, Grégoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel M. Kloumann, Ishan Misra, Ivan Evtimov, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Kartikeya Upasani, Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, and et al. 2024. [The llama 3 herd of models](https://arxiv.org/abs/2407.21783). *CoRR* abs/2407.21783. <https://doi.org/10.48550/ARXIV.2407.21783>.

OpenAI. 2023. [GPT-4 technical report](https://arxiv.org/abs/2303.08774). *CoRR* abs/2303.08774. <https://doi.org/10.48550/ARXIV.2303.08774>.